

Deep Reinforcement Learning

Prof. Joongheon Kim

Korea University, Seoul, Korea

<https://joongheon.github.io>

joongheon@korea.ac.kr

Lecture Roadmap

Introduction and Preliminaries

Deep Reinforcement Learning Theory

**Deep Reinforcement Learning
Implementation**

Imitation Learning and Autonomous Driving

- **Basics**
- Q-Learning Implementation
- DQN Implementation

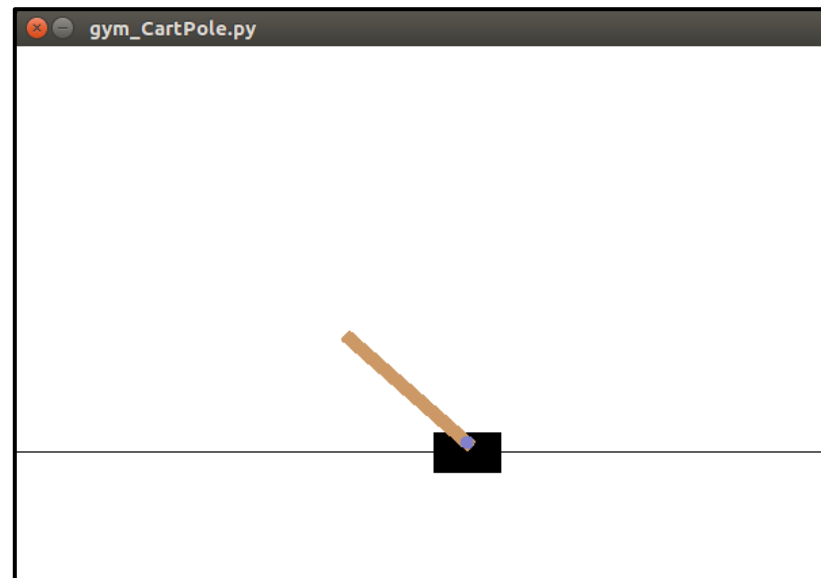
Deep Reinforcement Learning

DRL Implementation

- **Basics**
- Q-Learning Implementation
- DQN Implementation

Basics, Hello World: CartPole

```
1 import gym
2 env = gym.make('CartPole-v0')
3 env.reset()
4 for _ in range(1000):
5     env.render()
6     action = env.action_space.sample()
7     observation, reward, done, info = env.step(action)
8     #env.step(action)
```



Lecture Roadmap

Introduction and Preliminaries

Deep Reinforcement Learning Theory

**Deep Reinforcement Learning
Implementation**

- Basics
- **Q-Learning Implementation**
- DQN Implementation

DDPG-based Vehicular Caching

Imitation Learning and Autonomous Driving


Deep Reinforcement Learning

DRL Implementation

- Basics
- **Q-Learning Implementation**
- DQN Implementation


- Q-Learning Implementation
 - **Q-Learning (Basics)**
 - Q-Learning (Exploit and Exploration)

Q-Learning (Basics)



```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import gym
4 from gym.envs.registration import register
5 import random
6
7 '''
8 Q-Table
9 | action | L | D | R | U |
10 -----
11 state: 0 |   |   |   |   |
12 -----
13 state: 1 |   |   |   |   |
14 -----
15 state: 2 |   |   |   |   |
16 -----
17 state: ... |   |   |   |   |
18 -----
19 '''
20
21 register(
22     id='FrozenLake-v3',
23     entry_point='gym.envs.toy_text:FrozenLakeEnv',
24     kwargs={
25         'map_name': '4x4',
26         'is_slippery': False
27     }
28 )
29
30 env = gym.make("FrozenLake-v3")
```


Q-Learning (Basics)



```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import gym
4 from gym.envs.registration import register
5 import random
```

```
6
7 '''
8 Q-Table
9 |   | action | L | D | R | U |
10 -----
11 state: 0 |   |   |   |   |   |
12 -----
13 state: 1 |   |   |   |   |   |
14 -----
15 state: 2 |   |   |   |   |   |
16 -----
17 state: ... |   |   |   |   |   |
18 -----
19 '''
```

- Environment setting

```
21 register(
22     id='FrozenLake-v3',
23     entry_point='gym.envs.toy_text:FrozenLakeEnv',
24     kwargs={
25         'map_name': '4x4',
26         'is_slippery': False
27     }
28 )
29
30 env = gym.make("FrozenLake-v3")
```

Q-Learning (Basics)

```
32 # Initialization with 0 in Q-table
33 Q = np.zeros([env.observation_space.n, env.action_space.n]) # (16,4) where 16: 4*4 map, 4: actions
34 num_episodes = 1000 # Number of iterations
35
36 rList = []
37 successRate = []
38
39 def rargmax(vector):
40     m = np.amax(vector) # Return the maximum of an array or maximum along an axis (0 or 1)
41     indices = np.nonzero(vector == m)[0] # np.nonzero(True/False vector) => find the maximum
42     return random.choice(indices) # Random selection
43
44 for i in range(num_episodes): # Updates with num_episodes iterations
45     state = env.reset() # Reset
46     total_reward = 0 # Reward graph (1: success, 0: failure)
47     done = None
48
49     while not done: # The agent is not in the goal yet
50         action = rargmax(Q[state, :]) # Find maximum reward among 4 actions, find next action
51         new_state, reward, done, _ = env.step(action) # Result of the chosen action
52
53         Q[state, action] = reward + np.max(Q[new_state, :]) # Q-update
54         total_reward += reward
55         state = new_state
56
57     rList.append(total_reward) # Reward appending
58     successRate.append(sum(rList)/(i+1)) # Success rate appending
```

Q-Learning (Basics)

```
32 # Initialization with 0 in Q-table
33 Q = np.zeros([env.observation_space.n, env.action_space.n]) # (16,4) where 16: 4*4 map, 4: actions
34 num_episodes = 1000 # Number of iterations
35
36 rList = []
37 successRate = []
38
39 def rargmax(vector):
40     m = np.amax(vector) # Return the maximum of an array or maximum along an axis (0 or 1)
41     indices = np.nonzero(vector == m)[0] # np.nonzero(True/False vector) => find the maximum
42     return random.choice(indices) # Random selection
43
44 for i in range(num_episodes): # Updates with num_episodes iterations
45     state = env.reset() # Reset
46     total_reward = 0 # Reward graph (1: success, 0: failure)
47     done = None
48
49     while not done: # The agent is not in the goal yet
50         action = rargmax(Q[state, :]) # Find maximum reward among 4 actions, find next action
51         new_state, reward, done, _ = env.step(action) # Result of the chosen action
52
53         Q[state, action] = reward + np.max(Q[new_state, :]) # Q-update
54         total_reward += reward
55         state = new_state
56
57     rList.append(total_reward) # Reward appending
58     successRate.append(sum(rList)/(i+1)) # Success rate appending
```

- Randomly pick one when multiple argmax values exist


Q-Learning (Basics)

```
32 # Initialization with 0 in Q-table
33 Q = np.zeros([env.observation_space.n, env.action_space.n]) # (16,4) where 16: 4*4 map, 4: actions
34 num_episodes = 1000 # Number of iterations
35
36 rList = []
37 successRate = []
38
39 def rargmax(vector):
40     m = np.amax(vector) # Return the max
41     indices = np.nonzero(vector == m)[0]
42     return random.choice(indices) # Random choice
43
44 for i in range(num_episodes): # Updates
45     state = env.reset() # Reset
46     total_reward = 0 # Reward graph (1:
47     done = None
48
49     while not done: # The agent is not in the goal yet
50         action = rargmax(Q[state, :]) # Find maximum reward among 4 actions, find next action
51         new_state, reward, done, _ = env.step(action) # Result of the chosen action
52
53         Q[state, action] = reward + np.max(Q[new_state, :]) # Q-update
54         total_reward += reward
55         state = new_state
56
57     rList.append(total_reward) # Reward appending
58     successRate.append(sum(rList)/(i+1)) # Success rate appending
```

- Iteration until the agent arrives at the goal or it cannot move anymore.
- (line 50) find the action which returns max Q value.
- (line 51) take the action which is the result of (line 50).
- done: if the agent is at goal or cannot move anymore, done → True
- (line 53) Q-update
- (line 54) reward value accumulation
- (line 55) state value update for next iteration

Q-Learning (Basics)

```
68 '''
69 Final Q-Table [L, D, R, U]
70 [[0. 1. 0. 0.] 0 (D)
71 [0. 0. 0. 0.] 1
72 [0. 0. 0. 0.] 2
73 [0. 0. 0. 0.] 3
74 [0. 1. 0. 0.] 4 (D)
75 [0. 0. 0. 0.] 5
76 [0. 0. 0. 0.] 6
77 [0. 0. 0. 0.] 7
78 [0. 0. 1. 0.] 8 (R)
79 [0. 1. 0. 0.] 9 (D)
80 [0. 1. 0. 0.] 10 (D)
81 [0. 0. 0. 0.] 11
82 [0. 0. 0. 0.] 12
83 [0. 0. 1. 0.] 13 (R)
84 [0. 0. 1. 0.] 14 (R)
85 [0. 0. 0. 0.]] 15
86 Success Rate : 0.903
87 '''
```



- Q-Learning Implementation
 - Q-Learning (Basics)
 - **Q-Learning (Exploit and Exploration)**

Q-Learning (Exploit and Exploration)

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import gym
4 from gym.envs.registration import register
5 import random
6
7 register(
8     id='FrozenLake-v3',
9     entry_point='gym.envs.toy_text:FrozenLakeEnv',
10     kwargs={
11         'map_name': '4x4',
12         'is_slippery': False
13     }
14 )
15
16 env = gym.make("FrozenLake-v3")
17
18 Q = np.zeros([env.observation_space.n, env.action_space.n])
19 num_episodes = 1000
20
21 rList = []
22 successRate = []
23 e = 0.1 # exploit and exploration
24
25 mode = input(
26     "Mode Selection [(1) e-greedy (2) decaying e-greedy (3) random noise (etc) original]: ")
27 r = 0.9 # discount factor
```

Q-Learning (Exploit and Exploration)

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import gym
4 from gym.envs.registration import register
5 import random
6
7 register(
8     id='FrozenLake-v3',
9     entry_point='gym.envs.toy_text:FrozenLakeEnv',
10     kwargs={
11         'map_name': '4x4',
12         'is_slippery': False
13     }
14 )
15
16 env = gym.make("FrozenLake-v3")
17
18 Q = np.zeros([env.observation_space.n, env.action_space.n])
19 num_episodes = 1000
20
21 rList = []
22 successRate = []
23 e = 0.1 # exploit and exploration
24
25 mode = input(
26     "Mode Selection [(1) e-greedy (2) decaying e-greedy (3) random noise (etc) original]: ")
27 r = 0.9 # discount factor
```

- Parameter setting

Q-Learning (Exploit and Exploration)

```
29 def rargmax(vector):
30     m = np.amax(vector)
31     indices = np.nonzero(vector == m)[0]
32     return random.choice(indices)
33
34 for i in range(num_episodes):
35     state = env.reset()
36     total_reward = 0
37     done = None
38
39     while not done:
40         rand = random.random()
41         # e-greedy / decaying e-greedy
42         if (mode == '1' and rand < e) or (mode == '2' and (rand < e / (i + 1))):
43             action = env.action_space.sample()
44             # random noise
45         elif mode == '3':
46             action = rargmax(
47                 Q[state, :] + np.random.random(env.action_space.n) / (i + 1))
48             # original
49         else:
50             action = rargmax(Q[state, :])
51
52         new_state, reward, done, _ = env.step(action)
53         Q[state, action] = reward + r * np.max(Q[new_state, :])
54         total_reward += reward
55         state = new_state
56
57     rList.append(total_reward)
58     successRate.append(sum(rList) / (i + 1))
```

Q-Learning (Exploit and Exploration)

```
60 print("Final Q-Table")
61 print(Q)
62 print("Success Rate : ", successRate[-1])
63 plt.plot(range(len(rList)), rList)
64 plt.plot(range(len(successRate)), successRate)
65 plt.show()
66 '''
67 Mode Selection [(1) e-greedy (2) decaying e-greedy (3) random noise (etc) original]: 2
68 Final Q-Table
69 [[0.      0.      0.59049 0.      ]
70  [0.      0.      0.6561 0.      ]
71  [0.59049 0.729   0.      0.      ]
72  [0.      0.      0.      0.      ]
73  [0.      0.      0.      0.      ]
74  [0.      0.      0.      0.      ]
75  [0.      0.81    0.      0.      ]
76  [0.      0.      0.      0.      ]
77  [0.      0.      0.      0.      ]
78  [0.      0.      0.81    0.      ]
79  [0.      0.9     0.      0.      ]
80  [0.      0.      0.      0.      ]
81  [0.      0.      0.      0.      ]
82  [0.      0.      0.      0.      ]
83  [0.      0.      1.      0.      ]
84  [0.      0.      0.      0.      ]]
85 Success Rate : 0.932
86 '''
```


Q-Learning (Exploit and Exploration)

[L, D, R, U]

Final Q-Table

[0.	0.	0.59049	0.]	0 (R)
[0.	0.	0.6561	0.]	1 (R)
[0.59049	0.729	0.	0.]	2 (D)
[0.	0.	0.	0.]	3
[0.	0.	0.	0.]	4
[0.	0.	0.	0.]	5
[0.	0.81	0.	0.]	6 (D)
[0.	0.	0.	0.]	7
[0.	0.	0.	0.]	8
[0.	0.	0.81	0.]	9 (R)
[0.	0.9	0.	0.]	10 (D)
[0.	0.	0.	0.]	11
[0.	0.	0.	0.]	12
[0.	0.	0.	0.]	13
[0.	0.	1.	0.]	14 (R)
[0.	0.	0.	0.]]	15

Success Rate : 0.932



Lecture Roadmap

Introduction and Preliminaries

Deep Reinforcement Learning Theory

**Deep Reinforcement Learning
Implementation**

- Basics
- Q-Learning Implementation
- **DQN Implementation**

DDPG-based Vehicular Caching

Imitation Learning and Autonomous Driving


Deep Reinforcement Learning

DRL Implementation

- Basics
- Q-Learning Implementation
- **DQN Implementation**

- Frozen Lake

- Input: States 0~15 (totally 16) → one-hot encoding
- Output: 4 actions (totally 4) → Q-values for Up, Down, Left, and Right



```
1 import tensorflow as tf
2 import numpy as np
3
4 class DQN:
5     def __init__(self, session, input_size, output_size, name="main"):
6         self.session = session
7         self.input_size = input_size
8         self.output_size = output_size
9         self.net_name = name
10        self._build_network()
11
12    # 네트워크 구성 (레이어, 활성화 함수)
13    def _build_network(self, h_size=6, l_rate=1e-1):
14        with tf.variable_scope(self.net_name):
15            self._X = tf.placeholder(tf.float32, [None, self.input_size], name="input_x")
16
17            W1 = tf.get_variable("W1", shape=[
18                self.input_size, h_size], initializer=tf.contrib.layers.xavier_initializer())
19            layer1 = tf.nn.tanh(tf.matmul(self._X, W1))
20
21            W2 = tf.get_variable("W2", shape=[
22                h_size, self.output_size], initializer=tf.contrib.layers.xavier_initializer())
23            self._Qpredict = tf.matmul(layer1, W2)
24
25            self._Y = tf.placeholder(shape=[None, self.output_size], dtype=tf.float32)
26            self._loss = tf.reduce_mean(tf.square(self._Y - self._Qpredict))
27            self._train = tf.train.AdamOptimizer(learning_rate=l_rate).minimize(self._loss)
28
29    # x에 대한 y 결과 리턴 (2단계 네트워크 지나는 결과) -> 관계식 적용 결과
30    def predict(self, state):
31        x = np.reshape(state, [1, self.input_size])
32        return self.session.run(self._Qpredict, feed_dict={self._X:x})
33
34    # 들어온 데이터를 바탕으로 W 업데이트 (학습 시킨다)
35    def update(self, x_stack, y_stack):
36        return self.session.run([self._loss, self._train], feed_dict={self._X:x_stack, self._Y:y_stack})
```

DQN (NIPS 2013)

File Name:
dqn_2013.py

```
1  import numpy as np
2  import gym
3  from gym.envs.registration import register
4  import random
5  import tensorflow as tf
6  import class_dqn
7  from collections import deque
8  import matplotlib.pyplot as plt
9
10 env = gym.make("CartPole-v1")
11 env.max_episode_steps = 500
12
13 input_size = env.observation_space.shape[0]
14 output_size = env.action_space.n
15
16 learning_rate = 1e-1
17 discount_rate = 0.9
18 REPLAY_MEMORY = 50000
19 results = []
20
21 '''
22 DQN에는 두가지 문제가 있어 결과가 최상으로 나오지 않는다.
23 1) Correlation between samples
24 2) Non-stationary Targets
25 이 효과를 상쇄하기 위해서 2가지 방법을 사용하는데
26 하나는 레이어를 하나 더 만드는 것이고 또 하나는 버퍼에서 샘플링을 하여 학습에 사용하는 것이다.
27 '''
```



```
29 def simple_replay_train(DQN, train_batch):
30     # Array of uninitialized (arbitrary) data of the given shape
31     x_stack = np.empty(0).reshape(0, DQN.input_size)
32     y_stack = np.empty(0).reshape(0, DQN.output_size)
33
34     for state, action, reward, next_state, done in train_batch:
35         # 현재 가지고 있는 W1, W2로 예측한 Q값
36         Q = DQN.predict(state)
37
38         if done:
39             Q[0, action] = reward
40         else:
41             Q[0, action] = reward + discount_rate * \
42                 np.max(DQN.predict(next_state))
43
44         # 트레이닝할 Data 만들기
45         # X 값 : state, Y 값 : Q값 배열
46         x_stack = np.vstack([x_stack, state])
47         y_stack = np.vstack([y_stack, Q])
48
49     # 금방 들어온 데이터로 학습시키기
50     return DQN.update(x_stack, y_stack)
51
52 def bot_play(main_dqn):
53     s = env.reset()
54     reward_sum = 0
55     while True:
56         env.render()
57         a = np.argmax(main_dqn.predict(s))
58         s, reward, done, _ = env.step(a)
59         reward_sum += reward
60         if done:
61             print("Total score : {}".format(reward_sum))
62             break
```

DQN (NIPS 2013)

```
64 def main():
65     max_episodes = 2000
66     replay_buffer = deque() # 표본을 저장할 버퍼, 이곳에서 샘플링을 하여 학습에 사용한다.
67
68     with tf.Session() as sess:
69         main_dqn = class_dqn.DQN(sess, input_size, output_size, name="main") # DQN 클래스는 Q-Network의 속성과 멤버함수가 작성되어 있다. ( dqn.py )
70         tf.global_variables_initializer().run()
71
72         for i in range(max_episodes):
73             state = env.reset()
74             e = 1. / ((i / 10) + 1)
75             step_count = 0 # 한 번 테스트에 최대 몇번까지 움직였는가
76             done = False
77
78             while not done:
79                 # e-greedy 방식으로 action 선택
80                 if np.random.rand(1) < e:
81                     action = env.action_space.sample()
82                 else:
83                     # 현재 가지고 있는 W로 Q predict
84                     action = np.argmax(main_dqn.predict(state)) # 현재 state에서 현재 가지고 있는 W, b 값으로 Q 계산 후 action 선택
85
86                 next_state, reward, done, _ = env.step(action)
87
88                 # 안 끝나는 게 좋은 거니까 끝나면 reward -10; 그리고 학습은 시키지 않아
89                 if done:
90                     reward = -10
91
92                 # 현재 가지고 있는 W 값으로 얻은 결과를 쌓음
93                 replay_buffer.append((state, action, reward, next_state, done)) # 버퍼에 input과 output을 같이 저장
94                 if len(replay_buffer) > REPLAY_MEMORY: # 50000개까지만 저장하고 남으면 오래된 데이터 삭제
95                     replay_buffer.popleft()
96
```

DQN (NIPS 2013)

```
97         state = next_state
98         step_count += 1
99         # 너무 오래 하지 않도록 중간에 푸르닝
100         if step_count > 10000:
101             break
102
103     print("Episode : {}, steps : {}".format(i, step_count))
104     results.append(step_count)
105
106     if i % 10 == 1: # 10번 테스트 중 1번만 학습
107         for _ in range(50):
108             minibatch = random.sample(replay_buffer, 10) # 버퍼에서 10개를 샘플링
109             loss, _ = simple_replay_train(main_dqn, minibatch) # 샘플링한 10개를 바탕으로 학습
110             print("Loss : ", loss)
111
112     bot_play(main_dqn) # 최종 결과로 마지막 게임 수행
113     plt.title("(2013) Total step count on each episode")
114     plt.plot(range(len(results)), results)
115     plt.show()
116
117 if __name__ == "__main__":
118     main()
119
```